

Reducing Conflict Misses with Set-Associative Caches

Not too conflict-y. Not too slow.
... Just Right!



8 byte, 2-way set associative Cache

XXXX

CACHE

index	V	tag	data
0	0	xx	E F
1	0	xx	C D

MEMORY

addr	data
0000	A
0001	B
0010	C
0011	D
0100	E
0101	F
0110	G
0111	H
1000	J
1001	K
1010	L
1011	M
1100	N
1101	O
1110	P
1111	Q

What should the **offset** be?

What should the **index** be?

What should the **tag** be?

8 byte, 2-way set associative Cache

XXXX
tag | index | offset

CACHE

index	V	tag	data
0	0	xx	X X
1	0	xx	X X

load 0x1100 Miss
load 0x1101
load 0x0100
load 0x1100

Lookup:
⇒ Index into \$
⇒ Check tag
⇒ Check valid bit

MEMORY

addr	data
0000	A
0001	B
0010	C
0011	D
0100	E
0101	F
0110	G
0111	H
1000	J
1001	K
1010	L
1011	M
1100	N
1101	O
1110	P
1111	Q

8 byte, 2-way set associative Cache

XXXX
tag | index | offset

CACHE

index	V	tag	data
0	1	11	N O
1	0	xx	X X

load 0x1100 Miss
load 0x1101 Hit!
load 0x0100
load 0x1100

Lookup:
⇒ Index into \$
⇒ Check tag
⇒ Check valid bit

MEMORY

addr	data
0000	A
0001	B
0010	C
0011	D
0100	E
0101	F
0110	G
0111	H
1000	J
1001	K
1010	L
1011	M
1100	N
1101	O
1110	P
1111	Q

8 byte, 2-way set associative Cache

XXXX
tag | index | offset

CACHE

index	V	tag	data
0	1	11	N O
1	0	xx	X X

load 0x1100 Miss
load 0x1101 Hit!
load 0x0100 Miss
load 0x1100

Lookup:
⇒ Index into \$
⇒ Check tag
⇒ Check valid bit

MEMORY

addr	data
0000	A
0001	B
0010	C
0011	D
0100	E
0101	F
0110	G
0111	H
1000	J
1001	K
1010	L
1011	M
1100	N
1101	O
1110	P
1111	Q

8 byte, 2-way set associative Cache

XXXX
tag | index | offset

CACHE

index	V	tag	data
0	1	11	N O
1	0	xx	X X

load 0x1100 Miss
load 0x1101 Hit!
load 0x0100 Miss
load 0x1100 Hit!

Lookup:
⇒ Index into \$
⇒ Check tag
⇒ Check valid bit

MEMORY

addr	data
0000	A
0001	B
0010	C
0011	D
0100	E
0101	F
0110	G
0111	H
1000	J
1001	K
1010	L
1011	M
1100	N
1101	O
1110	P
1111	Q

Misses: the Three C's



Cold (compulsory) Miss:
never seen this address before



Conflict Miss:
cache associativity is too low



Capacity Miss:
cache is too small

ABCs of Caches



$$t_{avg} = t_{hit} + \%_{miss} * t_{miss}$$

+ Associativity:

↓ conflict misses ☺

↑ hit time ☹

+ Block Size:

↓ cold misses ☺

↑ conflict misses ☹


+ Capacity:

↓ capacity misses ☺

↑ hit time ☹

Which caches get what properties?

$$t_{avg} = t_{hit} + \%_{miss} * t_{miss}$$




Summary so far

- Things we've covered:
 - The Need for Speed
 - Locality to the Rescue!
 - Calculating average memory access time
 - \$ Misses: Cold, Conflict, Capacity
 - \$ Characteristics: Associativity, Block Size, Capacity
- Things we skipped (and are about to cover):
 - Cache Overhead
 - Replacement Policies
 - Writes

Basic Memory Array Structure

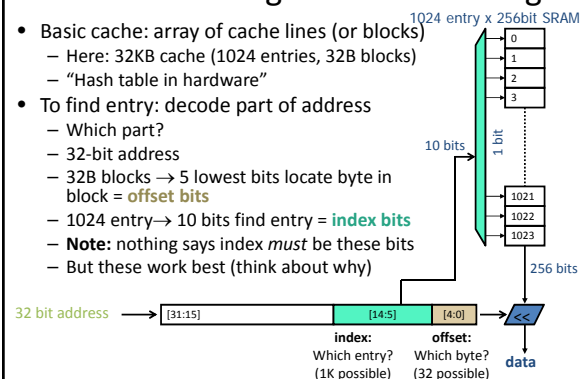
1024 entry x 256bit SRAM

- Number of entries
 - n bits for lookup $\rightarrow 2^n$ entries
 - Example: 1024 entries, 10 bit address
 - Decoder changes n-bit address to 2^n bit "one-hot" signal
- Size of entries
 - Width of data accessed
 - Here: 256 bits (32 bytes)



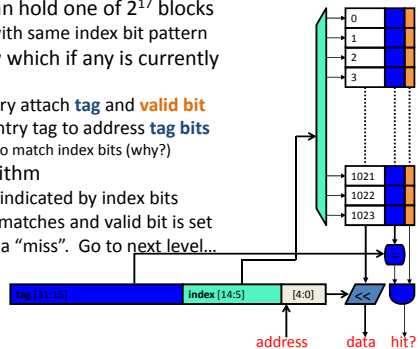
Caches: Finding Data via Indexing

- Basic cache: array of cache lines (or blocks)
 - Here: 32KB cache (1024 entries, 32B blocks)
 - "Hash table in hardware"
- To find entry: decode part of address
 - Which part?
 - 32-bit address
 - 32B blocks \rightarrow 5 lowest bits locate byte in block = **offset bits**
 - 1024 entry \rightarrow 10 bits find entry = **index bits**
 - **Note:** nothing says index *must* be these bits
 - But these work best (think about why)



Knowing that You Found It: Tags

- Each entry can hold one of 2^{17} blocks
 - All blocks with same index bit pattern
- How to know which if any is currently there?
 - To each entry attach **tag** and **valid bit**
 - Compare entry tag to address **tag bits**
 - No need to match index bits (why?)
- Lookup algorithm
 - Read entry indicated by index bits
 - "Hit" if tag matches and valid bit is set
 - Otherwise, a "miss". Go to next level...



Calculating Tag Overhead

- "32KB cache" means cache holds 32KB of **data**
 - Called **capacity**
 - Tag storage is considered overhead
- Tag overhead of 32KB cache with 1024 x 32B entries
 - 32B blocks → ??-bit offset
 - 1024 entries → ??-bit index
 - 32-bit address → ??-bit tag
- What about 64-bit addresses?

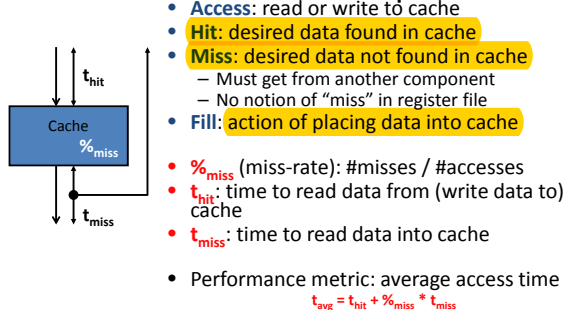
Calculating Tag Overhead

- "32KB cache" means cache holds 32KB of **data**
 - Called **capacity**
 - Tag storage is considered overhead
- Tag overhead of 32KB cache with 1024 x 32B entries
 - 32B blocks → 5-bit offset
 - 1024 entries → 10-bit index
 - 32-bit address → 32-bits – (5-bit offset + 10-bit index) = 17-bit tag
(17-bit tag + 1-bit valid) X 1024 entries = 18Kb tags = 2.2KB tags
~6% overhead
- What about 64-bit addresses?
 - Tag increases to 49 bits, ~20% overhead

Handling a Cache Miss

- What if requested data isn't in the cache?
 - How does it get in there?
- Cache controller**: finite state machine
 - Remembers miss address
 - Accesses next level of memory
 - Waits for response
 - Writes data/tag into proper locations
- All of this happens on the **fill path**
 - Sometimes called **backside**

Cache Performance Equation



CPI Calculation with Cache Misses

- Parameters**
 - Simple pipeline with base CPI of 1
 - Instruction mix: 30% loads/stores
 - IS: $\%_{miss} = 2\%$, $t_{miss} = 10$ cycles
 - DS: $\%_{miss} = 10\%$, $t_{miss} = 10$ cycles
- What is new CPI?**
 - $CPI_{IS} =$
 - $CPI_{DS} =$
 - $CPI_{new} =$

CPI Calculation with Cache Misses

Parameters

- Simple pipeline with base CPI of 1
- Instruction mix: 30% loads/stores
- IS: $\%_{\text{miss}} = 2\%$, $t_{\text{miss}} = 10$ cycles
- DS: $\%_{\text{miss}} = 10\%$, $t_{\text{miss}} = 10$ cycles

What is new CPI?

- $\text{CPI}_{\text{IS}} = \%_{\text{missIS}} * t_{\text{miss}} = 0.02 * 10 \text{ cycles} = 0.2 \text{ cycle}$
- $\text{CPI}_{\text{DS}} = \%_{\text{load/store}} * \%_{\text{missDS}} * t_{\text{missDS}} = 0.3 * 0.1 * 10 \text{ cycles} = 0.3 \text{ cycle}$
- $\text{CPI}_{\text{new}} = \text{CPI} + \text{CPI}_{\text{IS}} + \text{CPI}_{\text{DS}} = 1 + 0.2 + 0.3 = 1.5$

Measuring Cache Performance

- Ultimate metric is t_{avg}
 - Cache capacity and circuits roughly determines t_{hit}
 - Lower-level memory structures determine t_{miss}
- Measure $\%_{\text{miss}}$
 - Hardware performance counters
 - Simulation
 - Paper simulation (next)

Cache Miss Paper Simulation

- 4-bit addresses \rightarrow total memory size = 16B
 - Simpler cache diagrams than 32-bits
- 8B cache, 2B blocks
 - Number of entries (or sets): 4 (capacity / block-size)
 - Figure out how address splits into offset/index/tag bits
 - Offset:** least-significant $\log_2(\text{block-size}) = \log_2(2) = 1 \rightarrow 0000$
 - Index:** next $\log_2(\text{number-of-entries}) = \log_2(4) = 2 \rightarrow 0000$
 - Tag:** rest = $4 - 1 - 2 = 1 \rightarrow 0000$
- Cache diagram
 - 0000 | 0001 = addresses of data in block, values don't matter

Cache contents				Address	Outcome
Set00	Set01	Set10	Set11		
0000 0001	0010 0011	0100 0101	0110 0111		

Cache Miss Paper Simulation

8B cache, 2B blocks

Cache contents (prior to access)				Address	Outcome
Set00	Set01	Set10	Set11		
0000 0001	0010 0011	0100 0101	0110 0111	1100	Miss
				1000	
				1110	
				1000	
				0000	
				1000	

- How to reduce $\%_{\text{miss}}$? And hopefully t_{avg} ?

Cache Miss Paper Simulation


8B cache, 2B blocks

Cache contents (prior to access)				Address	Outcome
Set00	Set01	Set10	Set11		
0000 0001	0010 0011	0100 0101	0110 0111	1100	Miss
0000 0001	0010 0011	1100 1101	0110 0111	1110	Miss
0000 0001	0010 0011	1100 1101	1110 1111	1000	Miss
1000 1001	0010 0011	1100 1101	1110 1111	0011	Hit
1000 1001	0010 0011	1100 1101	1110 1111	1000	Hit
1000 1001	0010 0011	1100 1101	1110 1111	0000	Miss
0000 0001	0010 0011	1100 1101	1110 1111	1000	Miss

- How to reduce $\%_{\text{miss}}$? And hopefully t_{avg} ?

Capacity and Performance


- Simplest way to reduce $\%_{\text{miss}}$: increase capacity
 - Miss rate decreases monotonically
 - "Working set": insns/data program is actively using
 - Diminishing returns
 - However t_{hit} increases
 - Latency proportional to $\sqrt{\text{capacity}}$
- t_{avg} ?



- Given capacity, manipulate $\%_{\text{miss}}$ by changing **organization**

Block Size

- Given capacity, manipulate %_{miss} by changing organization
- One option: increase **block size**
 - Exploit **spatial locality**
 - Notice index/offset bits change
 - Tag remain the same
- Ramifications
 - + Reduce %_{miss} (up to a point)
 - + Reduce tag overhead (why?)
 - Potentially useless data transfer
 - Premature replacement of useful data
 - Fragmentation



Block Size and Tag Overhead

- Tag overhead of 32KB cache with 1024 32B entries
 - 32B lines → 5-bit offset
 - 1024 entries → 10-bit index
 - 32-bit address →
- Tag overhead of 32KB cache with 512 64B entries
 - 64B lines →
 - 512 entries →
 - 32-bit address →

Block Size and Tag Overhead

- Tag overhead of 32KB cache with 1024 32B entries
 - 32B lines → 5-bit offset
 - 1024 entries → 10-bit index
 - 32-bit address → 32 - (5-bit offset + 10-bit index) = 17-bit tag
 - (17-bit tag + 1-bit valid) X 1024 entries = 18Kb tags = 2.2KB tags
 - ~6% overhead
- Tag overhead of 32KB cache with 512 64B entries
 - 64B lines → 6-bit offset
 - 512 entries → 9-bit index
 - 32-bit address → 32 - (6-bit offset + 9-bit index) = 17-bit tag
 - (17-bit tag + 1-bit valid) X 512 entries = 9Kb tags = 1.1KB tags
 - ~3% overhead

Block Size Cache Miss Paper Simulation


8B cache, 4B blocks

		tag (1 bit)	index (1 bit)	offset (2 bits)	
Cache contents (prior to access)					
Set0	Set1				
0000 0001 0010 0011	0100 0101 0110 0111	1100			Miss
0000 0001 0010 0011	1100 1101 1110 1111	1110			Hit (spatial locality)
0000 0001 0010 0011	1100 1101 1110 1111	1000			Miss
1000 1001 1010 1011	1100 1101 1110 1111	0011			Miss
0000 0001 0010 0011	1100 1101 1110 1111	1000			Miss
1000 1001 1010 1011	1100 1101 1110 1111	0000			Miss
0000 0001 0010 0011	1100 1101 1110 1111	1000			Miss

- + **Spatial "prefetching"**: miss on 1100 brought in 1110
- **Conflicts**: miss on 1000 kicked out 0011

Effect of Block Size on Miss Rate

- Two effects on miss rate
 - + **Spatial prefetching (good)**
 - For blocks with adjacent addresses
 - Turns miss/miss into miss/hit pairs
 - **Interference (bad)**
 - For blocks with non-adjacent addresses (but in adjacent entries)
 - Turns hits into misses by disallowing simultaneous residence
 - Consider entire cache as one big block
- Both effects always present
 - Spatial prefetching dominates initially
 - Depends on size of the cache
 - Good block size is 16-128B
 - Program dependent



Block Size and Miss Penalty

- Does increasing block size increase t_{miss} ?
 - Don't larger blocks take longer to read, transfer, and fill?
 - They do, but...
- t_{miss} of an isolated miss is not affected
 - Critical Word First / Early Restart (CRF/ER)**
 - Requested word fetched first, pipeline restarts immediately
 - Remaining words in block transferred/filled in the background
- t_{miss} 'es of a cluster of misses will suffer
 - Reads/transfers/fills of two misses can't happen at the same time
 - Latencies can start to pile up
 - This is a bandwidth problem (more later)

Conflicts

8B cache, 2B blocks

tag (1 bit) index (2 bits) 1 bit					
Cache contents (prior to access)				Address	Outcome
Set00	Set01	Set10	Set11		
0000 0001	0010 0011	0100 0101	0110 0111	1100	Miss
0000 0001	0010 0011	1100 1101	0110 0111	1110	Miss
0000 0001	0010 0011	1100 1101	1110 1111	1000	Miss
1000 1001	0010 0011	1100 1101	1110 1111	0011	Hit
1000 1001	0010 0011	1100 1101	1110 1111	1000	Hit
1000 1001	0010 0011	1100 1101	1110 1111	0000	Miss
0000 0001	0010 0011	1100 1101	1110 1111	1000	Miss

- Pairs like 0000/1000 **conflict**
 - Regardless of block-size (assuming capacity < 16)
 - Q: can we allow pairs like these to simultaneously reside?
 - A: yes, reorganize cache to do so

Set-Associativity

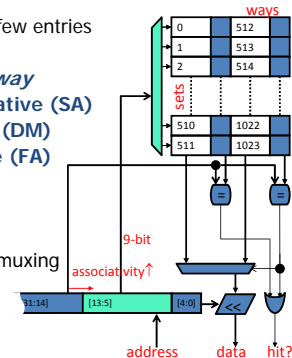
- Block can reside in one of few entries
- Entry groups called **sets**
- Each entry in set called a **way**
- This is **2-way set-associative (SA)**
- 1-way → **direct-mapped (DM)**
- 1-set → **fully-associative (FA)**

+ Reduces conflicts

- Increases latency_{hit}:

- additional tag match & muxing

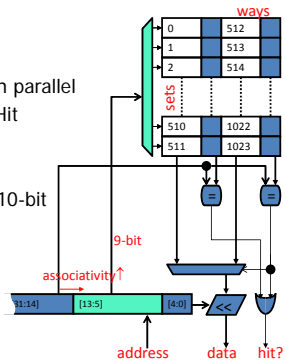
• Note: valid bit not shown



Set-Associativity

Lookup algorithm

- Use index bits to find set
- Read data/tags in all ways in parallel
- **Any** (match and valid bit), Hit
- Notice tag/index/offset bits
 - Only 9-bit index (versus 10-bit for direct mapped)
- Notice block numbering



Associativity and Miss Paper Simulation

8B cache, 2B blocks, **2-way set-associative**

Cache contents (prior to access)				Address	Outcome
Set0.Way0	Set0.Way1	Set1.Way0	Set1.Way1		
0000 0001	0100 0101	0010 0011	0110 0111	1100	Miss
				1110	
				1000	
				0011	
				1000	
				0000	
				1000	

+ **Avoid conflicts**: 0000 and 1000 can both be in set 0

- **Introduce some new conflicts**: addresses get re-arranged
 - Conflict avoidance usually dominates

Associativity and Miss Paper Simulation

8B cache, 2B blocks, **2-way set-associative**

Cache contents (prior to access)				Address	Outcome
Set0.Way0	Set0.Way1	Set1.Way0	Set1.Way1		
0000 0001	0100 0101	0010 0011	0110 0111	1100	Miss
1100 1101	0100 0101	0010 0011	0110 0111	1110	Miss
1100 1101	0100 0101	1110 1111	0110 0111	1000	Miss
1100 1101	1000 1001	1110 1111	0110 0111	0011	Miss (new conflict)
1100 1101	1000 1001	1110 1111	0010 0011	1000	Hit
1100 1101	1000 1001	1110 1111	0010 0011	0000	Miss
0000 0001	1000 1001	1110 1111	0010 0011	1000	Hit (avoid conflict)

+ **Avoid conflicts**: 0000 and 1000 can both be in set 0


- **Introduce some new conflicts**: addresses get re-arranged
 - Conflict avoidance usually dominates

Replacement Policies

- Associative caches present a new design choice
 - On cache miss, which block in set to replace (kick out)?
- Some options
 - **Random**
 - **FIFO (first-in first-out)**
 - **LRU (least recently used)**
 - Fits with temporal locality, LRU = least likely to be used in future
 - **NMRU (not most recently used)**
 - An easier to implement approximation of LRU
 - Is LRU for 2-way set-associative caches
 - **Belady's**: replace block that will be used furthest in future
 - Unachievable optimum
- Which policy is simulated in previous example?

Associativity and Performance

- Higher associative caches
 - + Have better (lower) $\%_{\text{miss}}$
 - Diminishing returns
 - However t_{hit} increases
 - The more associative, the slower
- What about t_{avg} ?



- Block-size and number of sets should be powers of two
 - Makes indexing easier (just rip bits out of the address)
- 3-way set-associativity? No problem