CSE 560
Computer Systems Architecture

Pipelining

Performance Review

What metric would you use to compare the performance of computers

1. With different ISAs?
2. With the same ISA?
3. With the same ISA and clock speed?
   - A. MIPS
   - B. Instructions/Program
   - C. Execution time
   - D. IPC
   - E. Clock speed

2

Performance Review

What metric would you use to compare the performance of computers

1. With different ISAs? **Execution time**
2. With the same ISA? **MIPS**
3. With the same ISA and clock speed? **IPC**
   - A. MIPS
   - B. Instructions/Program
   - C. Execution time
   - D. IPC
   - E. Clock speed

$$\frac{seconds}{program} = \frac{instructions}{program} \times \frac{cycles}{instruction} \times \frac{seconds}{cycle}$$

3

This Unit: (Scalar In-Order) Pipelining

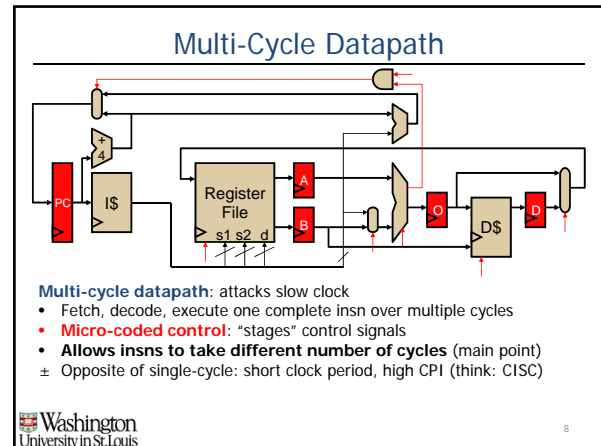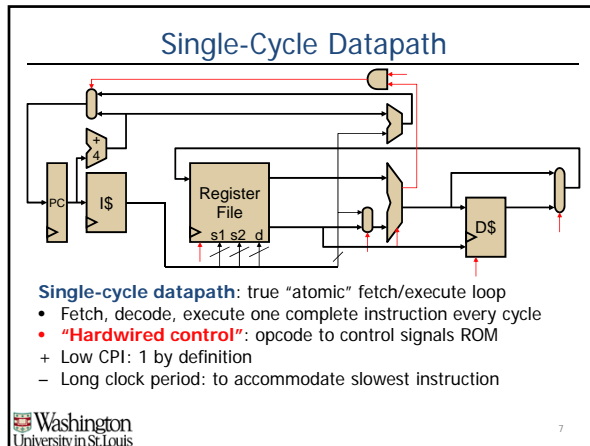| App	App	App
System software		
Mem	CPU	I/O

- Principles of pipelining
  - Effects of overhead and hazards
  - Pipeline diagrams
- Data hazards
  - Stalling and bypassing
- Control hazards (Next lecture)
  - Branch prediction
  - Predication

4

Datapath Background

5

Datapath and Control

Disclaimer: RISC datapath

- **Datapath**: implements execute portion of fetch/exec. loop
  - Functional units (ALUs), registers, memory interface
- **Control**: implements decode portion of fetch/execute loop
  - Mux selectors, write enable signals regulate flow of data in datapath
  - Part of decode involves translating insn opcode into control signals

6

Single-Cycle Datapath

Single-cycle datapath: true "atomic" fetch/execute loop
- Fetch, decode, execute one complete instruction every cycle
- **"Hardwired control"**: opcode to control signals ROM
- + Low CPI: 1 by definition
- – Long clock period: to accommodate slowest instruction

Multi-Cycle Datapath

Multi-cycle datapath: attacks slow clock
- Fetch, decode, execute one complete insn over multiple cycles
- **Micro-coded control**: "stages" control signals
- **Allows insns to take different number of cycles** (main point)
- ± Opposite of single-cycle: short clock period, high CPI (think: CISC)

Single-cycle vs. Multi-cycle Performance

- Single-cycle
  - Clock period = 50ns, CPI = 1
  - Performance = **50ns/insn**

- Multi-cycle has opposite performance split of single-cycle
  - + Shorter clock period
  - – Higher CPI

- Multi-cycle
  - Branch: 20% (**3** cycles), load: 20% (**5** cycles), ALU: 60% (**4** cycles)
  - Clock period = **11ns**, CPI = (20%*3)+(20%*5)+(60%*4) = 4
    - Why is clock period 11ns and not 10ns?
  - Performance = **44ns/insn**

- **Aside:** CISC makes perfect sense in multi-cycle datapath

Pipelining Basics

Latency versus Throughput

Single-cycle
Multi-cycle

- Can we have both low CPI and short clock period?
  - Not if datapath executes only one insn at a time
- Latency vs. Throughput
  - – Latency: no good way to make a single insn go faster
  - + **Throughput**: luckily, single insn latency not so important
    - Goal is to make programs, not individual insns, go faster
    - Programs contain billions of insns
 - Key: **exploit inter-insn parallelism**

Pipelining

Multi-cycle
Pipelined

- Important performance technique
  - **Improves insn throughput** rather instruction latency
- Begin with multi-cycle design
  - One insn advances from stage 1 to 2, next insn enters stage 1
  - Form of parallelism: "insn-stage parallelism"
  - Maintains illusion of sequential fetch/execute loop
  - Individual instruction takes the same number of stages
  - + **But instructions enter and leave at a much faster rate**
- Laundry analogy

Five Stage Pipelined Datapath

- Temporary values (PC,IR,A,B,O,D) re-latched every stage
  - Why? 5 insns may be in pipeline at once with different PCs
  - Notice, PC not latched after ALU stage (not needed later)
  - **Pipelined control**: one single-cycle controller
   - Control signals themselves pipelined

Five Stage Pipeline Performance

$T_{insn-mem}$ $T_{regfile}$ T_{ALU} $T_{data-mem}$ $T_{regfile}$

$T_{singlecycle}$

Pipelining: cut datapath into N stages (here five)
- One insn in each stage in each cycle
- + Clock period = $MAX(T_{insn-mem}, T_{regfile}, T_{ALU}, T_{data-mem})$
- + Base CPI = 1: insn enters and leaves every cycle
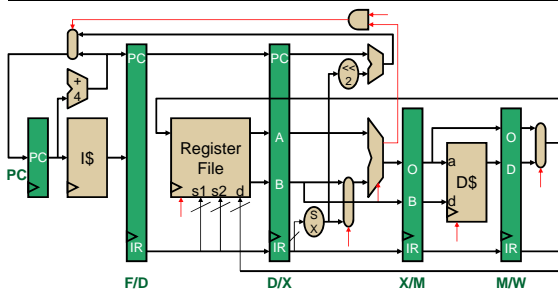- Individual insn latency increases (pipeline overhead), ok

Pipeline Terminology

PC   F/D   D/X   X/M   M/W

- Five stage: **F**etch, **D**ecode, e**X**ecute, **M**emory, **W**riteback
- Latches (pipeline registers) named by stages they separate
  - **PC**, **F/D**, **D/X**, **X/M**, **M/W**

More Terminology & Foreshadowing

- **Scalar pipeline**: one insn per stage per cycle
  - Alternative: "superscalar", *e.g.*, 4-wide (later)

- **In-order pipeline**: insns enter execute stage in order
  - Alternative: "out-of-order" (OoO) (later)

- **Pipeline depth**: number of pipeline stages
  - Nothing magical about five (Pentium 4 had 22 stages!)
  - Trend: deeper until Pentium 4, then pulled back a bit

Pipeline Example: Cycle 1

F/D   D/X   X/M   M/W

`add $3 $2,$1`

Pipeline Example: Cycle 2

F/D   D/X   X/M   M/W

`lw $4,0($5)`   `add $3 $2,$1`

Pipeline Example: Cycle 3

PC | I$ | Register File | D$

F/D   D/X   X/M   M/W

`sw $6,4($7)    lw $4,0($5)    add $3 $2,$1`
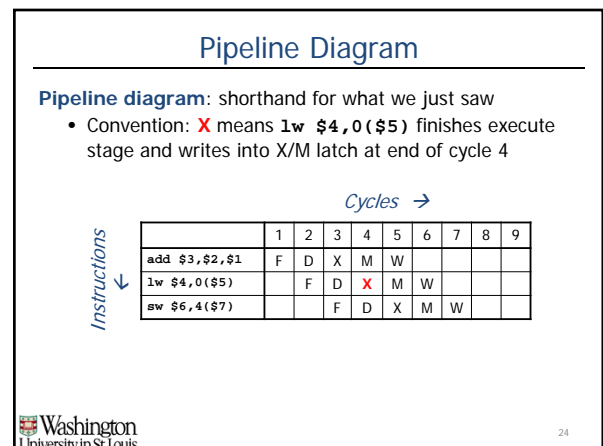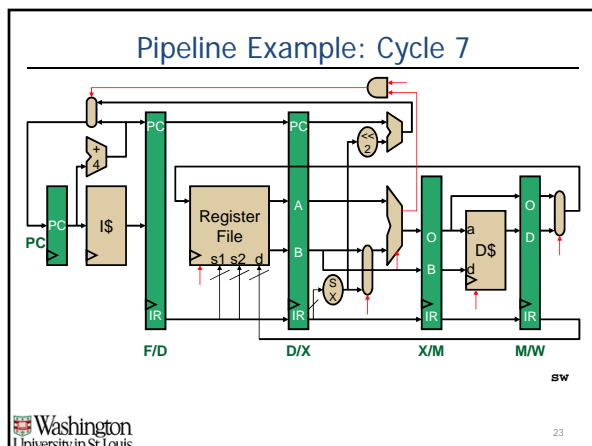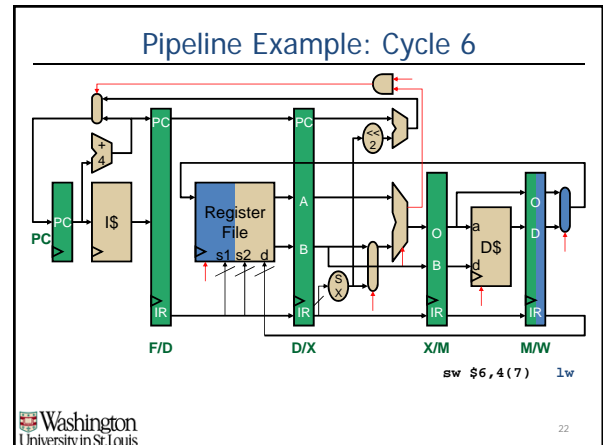
19

Pipeline Example: Cycle 4

PC | I$ | Register File | D$

F/D   D/X   X/M   M/W

`sw $6,4($7)    lw $4,0($5)    add $3 $2,$1`

20

Pipeline Example: Cycle 5

PC | I$ | Register File | D$

F/D   D/X   X/M   M/W

`sw $6,4($7)    lw $4,0($5)    add`

21

Pipeline Example: Cycle 6

PC | I$ | Register File | D$

F/D   D/X   X/M   M/W

`sw $6,4(7)    lw`

22

Pipeline Example: Cycle 7

PC | I$ | Register File | D$

F/D   D/X   X/M   M/W

`sw`

23

Pipeline Diagram

Pipeline diagram: shorthand for what we just saw
- Convention: **X** means `lw $4,0($5)` finishes execute stage and writes into X/M latch at end of cycle 4

Cycles →

Instructions ↓

	1	2	3	4	5	6	7	8	9
`add $3,$2,$1`	F	D	X	M	W				
`lw $4,0($5)`		F	D	**X**	M	W			
`sw $6,4($7)`			F	D	X	M	W		

24

Example Pipeline Perf. Calculation

- Single-cycle
  - Clock period = 50ns, CPI = 1
  - Performance = 50ns/insn
- Multi-cycle
  - Branch: 20% (3 cycles), load: 20% (5 cycles), ALU: 60% (4 cycles)
  - Clock period = 11ns, CPI = (20%*3)+(20%*5)+(60%*4) = 4
  - Performance = 44ns/insn
- 5-stage pipelined
  - Clock period = **12ns**   approx. (50ns / 5 stages) + overheads
  - + CPI = **1** (each insn takes 5 cycles, but 1 completes each cycle)
    - + Performance = **12ns/insn**
  - – Well actually ... CPI = 1 + some penalty for pipelining (next)
    - CPI = **1.5** (on average insn completes every 1.5 cycles)
    - Performance = **18ns/insn**
    - Much higher performance than single-cycle or multi-cycle

---

Clock Period of a Pipelined Processor

$Delay_{dp}$ = time it takes to travel through original datapath
N_{ps} = number of pipeline stages

Pipeline Clock Period > $Delay_{dp}$ / N_{ps}

- Latches add delay
- Extra "bypassing" logic adds delay
- Pipeline stages have different delays, clock period is max delay

- These factors have implications for ideal number pipeline stages
  - Diminishing clock frequency gains for longer (deeper) pipelines

---

CPI Calculation: Accounting for Stalls

Why is Pipelined CPI > 1 ?
- CPI for scalar in-order pipeline is 1 **+ stall penalties**
- Stalls used to resolve hazards
  - **Hazard**: condition that jeopardizes sequential illusion
  - **Stall**: pipeline delay introduced to restore sequential illusion
- Calculating pipeline CPI
  - **Frequency of stall** * **stall cycles**
  - Penalties add (stalls generally don't overlap in in-order pipelines)
  - 1 + stall-freq$_1$*stall-cyc$_1$ + stall-freq$_2$*stall-cyc$_2$ + ...
- Correctness/performance/make common case fast (MCCF)
  - Long penalties OK if rare, e.g., 1 + 0.01 * 10 = 1.1
  - Stalls have implications for ideal number of pipeline stages

---

Data Dependences, Pipeline Hazards, and Bypassing

Dependences and Hazards

- **Dependence**: relationship between two insns
  - **Data**: two insns use same storage location
  - **Control**: 1 insn affects whether another executes at all
  - *Not a bad thing*, programs would be boring otherwise
  - Enforced by making older insn go before younger one
    - Happens naturally in single-/multi-cycle designs
    - But not in a pipeline

- **Hazard**: dependence & possibility of wrong insn order
  - Effects of wrong insn order cannot be externally visible
    - **Stall**: for order by keeping younger insn in same stage
  - *Hazards are a bad thing*: stalls reduce performance

Why Does Every Insn Take 5 Cycles?

- Could/should we allow **add** to skip M and go to W?
  - It wouldn't help: peak fetch still only 1 insn per cycle
  - **Structural hazards**: who gets the register file write port?

Structural Hazards

- **Structural hazards**
  - Two insns trying to use same circuit at same time
    - E.g., structural hazard on register file write port

- **To fix structural hazards**: proper ISA/pipeline design
  - Each insn uses every structure exactly once
  - For at most one cycle
  - Always at same stage relative to F (fetch)

- **Tolerate structure hazards**
  - Add stall logic to stall pipeline when hazards occur

Example Structural Hazard

|	1	2	3	4	5	6	7	8	9
ld r2,0(r1)	F	D	X	**M**	W				
add r1 r3,r4		F	D	X	M	W			
sub r1 r3,r5			F	D	X	M	W		
st r6,0(r1)				**F**	D	X	M	W	

- **Structural hazard**: resource needed twice in one cycle
  - Example: unified instruction & data memories (caches)
  - Solutions:
    - Separate instruction/data memories (caches)
    - Have cache allow 2 accesses per cycle (slow, expensive)
    - Stall pipeline

Data Hazards

sw $3,0($7)   addi $6 1,$3   lw $4,0($3) add $3 $2,$1

- Would these instructions execute correctly on this pipeline?
- Which instructions execute with correct inputs?
  - **add** writes result into **$3** in current cycle
  - **lw** read **$3** two cycles ago → got wrong value
  - **addi** read **$3** one cycle ago → got wrong value
  - **sw** reads **$3** this cycle → maybe (depends on register file)

## Memory Data Hazards

lw $4,0($1)   sw $5,0($1)

- Are memory data hazards a problem for this pipeline? No
  - **lw** following **sw** to same address in next cycle, gets right value
  - Why? D$ read/write always take place in same stage
- Data hazards through registers? Yes (previous slide)
  - Occur because register write is three stages after register read
  - Can only read a register value three cycles after writing it

---

Observation!

lw $4,0($3)   add $3 $2,$1

- *Technically*, we have a problem:
  - **lw $4,0($3)** has already read **$3** from regfile
  - **add $3 $2,$1** hasn't yet written **$3** to regfile
- Fundamentally, this *should work*
  - **lw $4,0($3)** hasn't actually used **$3** yet
  - **add $3 $2,$1** has already computed **$3**

Reducing Data Hazards: Bypassing


lw $4,0($3)   add $3 $2,$1

Bypassing
- Reading a value from an intermediate (µarchitectural) source
- Not waiting until it is available from primary source
- Here, we bypass the register file
- Also called **forwarding**

WX Bypassing


```
Register
File
s1 s2  d
F      D          IR    X    IR    M    IR    W
                        lw $4,0($3)        add $3 $2,$1
```

- What about this combination?
  - Add another bypass path and MUX (multiplexor) input
  - First one was an **MX** bypass
 - This one is a **WX** bypass

---

ALUinB Bypassing


```
Register
File
s1 s2  d
F      D          IR    X    IR    M    IR    W
                        add $4 $2,$3       add $3 $2,$1
```

- Can also bypass to ALU input B

WM Bypassing?


```
Register
File
s1 s2  d
F      D          IR    X    IR    M    IR    W
                        sw $3,0($4)    lw $3,0($2)
```

- Does WM bypassing make sense?
  - Not to the address input (why not?)
  - But to the store data input, yes

Bypass Logic


```
Register
File
s1 s2  d
F      D          IR    X    IR    M    IR    W
                           bypass
```

Each MUX has its own logic; here it is for MUX ALUinA
  (D/X.IR.RegSource1 == X/M.IR.RegDest) => 0
  (D/X.IR.RegSource1 == M/W.IR.RegDest) => 1
  Else => 2


Pipeline Diagrams with Bypassing

- If bypass exists, "from"/"to" stages execute in same cycle
  - Example: full bypassing, use MX bypass

	1	2	3	4	5	6	7	8	9	10
add r2,r3➔r1	F	D	X	M	W					
sub r1,r4➔r2		F	D	X	M	W				

  - Example: full bypassing, use WX bypass

	1	2	3	4	5	6	7	8	9	10
add r2,r3➔r1	F	D	X	M	W					
ld [r7]➔r5		F	D	X	M	W				
sub r1,r4➔r2			F	D	X	M	W			

  - Example: WM bypass

	1	2	3	4	5	6	7	8	9	10
add r2,r3➔r1	F	D	X	M	W					
?		F	D	X	M	W				

Have We Prevented All Data Hazards?


```
Register
File
s1 s2  d
F      D  nop     IR    X    IR    M    IR    W
   stall
                        add $4 $2,$3       lw $3,4($2)
```

- No. Consider a "load" followed by a dependent "add" insn
- Bypassing alone isn't sufficient!
- Hardware solution: detect this situation and inject a stall cycle
- Software solution: ensure compiler doesn't generate such code

Stalling to Avoid Data Hazards

- Prevent F/D insn from reading (advancing) this cycle
  - Write **nop** into D/X.IR (effectively, insert **nop** in hardware)
  - Also reset (clear) the datapath control signals
  - Disable F/D latch and PC write enables (why?)
- Re-evaluate situation next cycle

Stalling on Load-To-Use Dependences


```
add $4 $2,$3      lw $3,4($2)
```
Stall = (D/X.IR.Operation == LOAD) &&
　　((F/D.IR.RegSrc1 == D/X.IR.RegDest) ||
　　((F/D.IR.RegSrc2 == D/X.IR.RegDest) && (F/D.IR.OP != STORE))

Stalling on Load-To-Use Dependences


```
add $4 $2,$3      (stall bubble)      lw $3,4($2)
```
Stall = (D/X.IR.Operation == LOAD) &&
　　((F/D.IR.RegSrc1 == D/X.IR.RegDest) ||
　　((F/D.IR.RegSrc2 == D/X.IR.RegDest) && (F/D.IR.OP != STORE))

Stalling on Load-To-Use Dependences


```
add $4 $2,$3      (stall bubble)      lw $3,…
```
Stall = (D/X.IR.Operation == LOAD) &&
　　((F/D.IR.RegSrc1 == D/X.IR.RegDest) ||
　　((F/D.IR.RegSrc2 == D/X.IR.RegDest) && (F/D.IR.OP != STORE))

Performance Impact of Load/Use Penalty

- Assume
  - Branch: 20%, load: 20%, store: 10%, other: 50%
  - 50% of loads are followed by dependent instruction
    - require 1 cycle stall (*i.e.*, insertion of 1 **nop**)

- Calculate CPI
  - CPI = 1 + (1 * 20% * 50%) = **1.1**

Reducing Load-Use Stall Frequency

	1	2	3	4	5	6	7	8	9
add $3 $2,$1	F	D	X	M	W				
lw $4,4($3)		F	D	X	M	W			
addi $6 $4,1			F	d*	D	X	M	W	
sub $8 $3,$1				F	D	X	M	W	

- Use compiler scheduling to reduce load-use stall frequency
  - More on compiler scheduling later

	1	2	3	4	5	6	7	8	9
add $3 $2,$1	F	D	X	M	W				
lw $4,4($3)		F	D	X	M	W			
sub $8 $3,$1			F	D	X	M	W		
addi $6 $4,1				F	D	X	M	W	

Pipelining and Multi-Cycle Operations

- What if you wanted to add a multi-cycle operation?
  - *E.g.*, 4-cycle multiply
  - **P/W**: separate output latch connects to W stage
  - Controlled by pipeline control finite state machine (FSM)

A Pipelined Multiplier

- Multiplier itself is often pipelined, what does this mean?
  - Product/multiplicand register/ALUs/latches replicated
  - Can start a new multiply operation every cycle

Pipeline Diagram with Multiplier

	1	2	3	4	5	6	7	8	9
mul $4 $3,$5	F	D	P0	P1	P2	P3	W		
addi $6 $4,1		F	d*	d*	d*	D	X	M	W

- What about...
  - Two instructions trying to write regfile in same cycle?
  - Structural hazard!
- Must prevent:

	1	2	3	4	5	6	7	8	9
mul $4 $3,$5	F	D	P0	P1	P2	P3	W		
addi $6 $1,1		F	D	X	M	W			
add $5 $6,$10			F	D	X	M	W		

More Multiplier Nasties

- What about...
  - Mis-ordered register writes
  - SW thinks `add` gets `$4` from `addi`, actually gets it from `mul`

	1	2	3	4	5	6	7	8	9
mul $4 $3,$5	F	D	P0	P1	P2	P3	W		
addi $4 $1,1		F	D	X	M	W			
…									
…									
add $10 $4,$6						F	D	X	M	W

- Common? Not for a 4-cycle multiply with 5-stage pipeline
  - More common with deeper pipelines
  - Frequency irrelevant: must be correct no matter how rare

Corrected Pipeline Diagram

- With the correct stall logic
  - Prevent mis-ordered writes to the same register
  - Why two cycles of delay?

	1	2	3	4	5	6	7	8	9
mul $4 $3,$5	F	D	P0	P1	P2	P3	W		
addi $4 $1,1		F	d*	d*	D	X	M	W	
…									
…									
add $10 $4,$6						F	D	X	M	W

Multi-cycle operations complicate pipeline logic

Pipelined Functional Units

- Almost all multi-cycle functional units are pipelined
  - Each operation takes N cycles
  - Can initiate a new (independent) operation every cycle
  - Requires internal latching and some hardware replication
  + Cheaper than multiple (non-pipelined) units

	1	2	3	4	5	6	7	8	9	10	11
mulf f0 f1,f2	F	D	E1	E2	E3	E4	W				
mulf f3 f4,f5		F	D	E1	E2	E3	E4	W			

- Exception: int/FP divide: difficult to pipeline; not worth it

	1	2	3	4	5	6	7	8	9	10	11
divf f0 f1,f2	F	D	E/	E/	E/	E/	W				
divf f3 f4,f5		F	s*	s*	s*	D	E/	E/	E/	E/	W

s* = structural hazard, two insns need same structure
  - ISAs and pipelines designed minimize these
  - Canonical example: all insns go through M stage

ISA Implementability Review

Give an example of an ISA feature that makes pipelining more difficult and the particular pipeline stages it affects.

ISA Implementability Review

Give an example of an ISA feature that makes pipelining more difficult and the particular pipeline stages it affects.

- Variable instruction length and format make pipelining fetch and decode difficult.
- Implicit state makes dynamic scheduling difficult.
- Variable latencies makes scheduling difficult.