CSE 560
Computer Systems Architecture

Static Scheduling

Multiple Issue Redux

Which of the following statements is false?

A. Multiple issue is needed to expose insn level parallelism (ILP) beyond pipelining
B. Multiple issue improves performance
C. Multiple issue improves utilization
D. Multiple issue doesn't make sense past 6-wide
E. Multiple issue can be a hardware or a software technique


Multiple Issue Redux

Which of the following statements is false?

A. Multiple issue is needed to expose insn level parallelism (ILP) beyond pipelining
B. Multiple issue improves performance
C. Multiple issue improves utilization
D. Multiple issue doesn't make sense past 6-wide
E. Multiple issue can be a hardware or a software technique

This Unit: Static Scheduling

App	App	App
System software		
Mem	CPU	I/O

- Pipelining and superscalar review
- Code scheduling to
  - Reduce pipeline stalls
  - Increase ILP

Two approaches to scheduling

- This Unit:
  - Static scheduling by the compiler
- Coming Soon:
  - Dynamic scheduling by the hardware

Code Scheduling

- **Scheduling:** act of finding independent instructions
  - **Static:** at compile time by the compiler (software)
  - **Dynamic:** at runtime by the processor (hardware)

- Why schedule code?
  - **Scalar pipelines:** fill load-to-use delays to improve CPI
  - **Superscalar:** place independent instructions together
    - As above, load-to-use delay slots
    - Allow multiple-issue decode logic to let them execute at the same time

---

Scheduling Requirements

- **Independent insns**
  - no ILP → game over
- **Large Scheduling Scope**
  - Scope = code region we are scheduling
  - The bigger the better (more independent insns to play with)
  - Once scope is defined, schedule is pretty obvious
  - Trick is creating a large scope (schedule across branches?)
- **Enough registers**
  - To hold additional "live" values
- **Alias analysis**
  - Whether load/store reference same memory locations
    - Can they be legally rearranged?

Scheduling Techniques

- **Stall Removal**
  - Separate load-use pairs

- **Scope enlarging**
  - For Loops: **loop unrolling**
  - For Non-loops:
    - Superblocks
    - Predication

- **Exploit Data-Level Parallelism**
  - Vectors

New Metric: Utilization

Utilization: actual performance / peak performance
- Important metric for performance/cost
- Why pay for hardware you rarely use?

- Adding hardware usually ↑ performance, ↓utilization
  - New hardware cannot always be exploited
  - Diminishing marginal returns

- Compiler can help make better use of existing hardware
  - Important for superscalar

Running Code Example: SAXPY

- **SAXPY** (Single-precision A X Plus Y)
  - Linear algebra routine (for solving systems of equations)
  - Part of early Livermore Loops benchmark suite
  - floating point uses "F" registers and "F" instructions

```
for (i=0;i<N;i++)
  Z[i]=(A*X[i])+Y[i];

0: ldf X(r1)➔f1      // loop              LOAD1
1: mulf f0,f1➔f2     // A in f0           USE1
2: ldf Y(r1)➔f3      // X,Y,Z constants   LOAD2
3: addf f2,f3➔f4                          USE2
4: stf f4➔Z(r1)
5: addi r1,4➔r1      // i in r1
6: blt r1,r2,0       // N*4 in r2
```

SAXPY Performance and Utilization

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
ldf X(r1)➔f1	F	D	X	M	W															
mulf f0,f1➔f2		F	D	d*	E*	E*	E*	E*	E*	W										
ldf Y(r1)➔f3			F	p*	D	X	M	W												
addf f2,f3➔f4				F	d*	d*	d*	D	E+	E+	W									
stf f4➔Z(r1)					F	p*	p*	p*	D	X	M	W								
addi r1,4➔r1									F	D	X	M	W							
blt r1,r2,0										F	D	X	M	W						
ldf X(r1)➔f1																F	D	X	M	W

Scalar pipeline
- Full bypassing, 5-cycle E*, 2-cycle E+, predict branches taken
- Single iteration (7 insns) latency: **16–5 = 11 cycles**
- **Performance**: 7 insns / 11 cycles = 0.64 IPC
- **Utilization**: actual/peak IPC = 0.64 / 1 = 64%

SAXPY Performance and Utilization

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
ldf X(r1)➔f1	F	D	X	M	W															
mulf f0,f1➔f2	F	D	d*	d*	E*	E*	E*	E*	W											
ldf Y(r1)➔f3		F	D	X	M	W														
addf f2,f3➔f4		F	p*	p*	d*	d*	d*	D	E+	E+	W									
stf f4➔Z(r1)			F	p*	p*	p*	p*	p*	D	X	M	W								
addi r1,4➔r1					F	p*	p*	p*	p*	p*	D	X	M	W						
blt r1,r2,0					F	p*	p*	p*	p*	p*	d*	D	X	M	W					
ldf X(r1)➔f1																F	D	X	M	W

2-way superscalar pipeline
- Any two insns per cycle + split integer and FP pipelines
- + **Performance**: 7 insns / 10 cycles = 0.70 IPC
- – **Utilization**: actual/peak IPC = 0.70 / 2 = 35%
- – More hazards → more stalls
- – Each stall is more expensive

Eliminate Load-Use Pairs?

```
for (i=0;i<N;i++)
  Z[i]=(A*X[i])+Y[i];
```

```
0: ldf X(r1)➔f1   LOAD1        0: ldf X(r1)➔f1   LOAD1
1: mulf f0,f1➔f2  USE1         2: ldf Y(r1)➔f3   LOAD2
2: ldf Y(r1)➔f3   LOAD2        1: mulf f0,f1➔f2  USE1
3: addf f2,f3➔f4  USE2         3: addf f2,f3➔f4  USE2
4: stf f4➔Z(r1)                4: stf f4➔Z(r1)
5: addi r1,4➔r1                5: addi r1,4➔r1
6: blt r1,r2,0                 6: blt r1,r2,0
```

Problem solved?

Loop Unrolling SAXPY

- **Goal**: separate dependent insns from one another
- SAXPY problem: not enough flexibility within one iteration
  - Longest chain of insns is 9 cycles
    - Load (1)
    - Forward to multiply (5)
    - Forward to add (2)
    - Forward to store (1)
  - Can't hide a 9-cycle chain using only 7 insns
  - But how about two 9-cycle chains using 14 insns?
- **Loop unrolling**: schedule 2+ iterations together
  - Fuse iterations
  - Schedule to reduce stalls
  - Schedule introduces ordering problems → rename registers

13

Unrolling SAXPY I: Fuse Iterations

- Combine two (in general K) iterations of loop
  - Fuse loop control: induction variable (`i=r1`) increment + branch
  - Adjust (implicit) induction uses: constants → constants + 4

```
ldf X(r1),f1                          ldf X(r1),f1
mulf f0,f1,f2                         mulf f0,f1,f2
ldf Y(r1),f3                          ldf Y(r1),f3
addf f2,f3,f4                         addf f2,f3,f4
stf f4,Z(r1)                         stf f4,Z(r1)
addi r1,4,r1  -- increment i
blt r1,r2,0   -- jump back
ldf X(r1),f1                          ldf X+4(r1),f1
mulf f0,f1,f2                         mulf f0,f1,f2
ldf Y(r1),f3                          ldf Y+4(r1),f3
addf f2,f3,f4                         addf f2,f3,f4
stf f4,Z(r1)                         stf f4,Z+4(r1)
addi r1,4,r1  -- increment i          addi r1,8,r1
blt r1,r2,0   -- jump back            blt r1,r2,0
```

14

Unrolling SAXPY II: Pipeline Schedule

- Pipeline schedule to reduce stalls
  - Have already seen this: pipeline scheduling

```
ldf X(r1),f1                          ldf X(r1),f1
mulf f0,f1,f2                         ldf X+4(r1),f1
ldf Y(r1),f3                          mulf f0,f1,f2
addf f2,f3,f4                         mulf f0,f1,f2
stf f4,Z(r1)                         ldf Y(r1),f3
ldf X+4(r1),f1                        ldf Y+4(r1),f3
mulf f0,f1,f2                         addf f2,f3,f4
ldf Y+4(r1),f3                        addf f2,f3,f4
addf f2,f3,f4                         stf f4,Z(r1)
stf f4,Z+4(r1)                       stf f4,Z+4(r1)
addi r1,8,r1                          addi r1,8,r1
blt r1,r2,0                           blt r1,r2,0
```

15

Unrolling SAXPY III: "Rename" Registers

- Pipeline scheduling causes reordering violations
  - Rename registers to correct

```
ldf X(r1),f1   problem!               ldf X(r1),f1
ldf X+4(r1),f1                        ldf X+4(r1),f5
mulf f0,f1,f2                         mulf f0,f1,f2
mulf f0,f1,f2                         mulf f0,f5,f6
ldf Y(r1),f3                         ldf Y(r1),f3
ldf Y+4(r1),f3                        ldf Y+4(r1),f7
addf f2,f3,f4                         addf f2,f3,f4
addf f2,f3,f4                         addf f6,f7,f8
stf f4,Z(r1)                         stf f4,Z(r1)
stf f4,Z+4(r1)                       stf f8,Z+4(r1)
addi r1,8,r1                          addi r1,8,r1
blt r1,r2,0                           blt r1,r2,0
```

Do we have enough registers to do this?

Are we sure we can move these loads above these stores?
Alias analysis must be conservative.

16

Unrolled SAXPY Performance/Utilization

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
ldf X(r1)➔f1	F	D	X	M	**W**															
ldf X+4(r1)➔f5		F	D	X	M	W														
mulf f0,f1➔f2			F	D	E*	E*	E*	E*	**E***	W										
mulf f0,f5➔f6				F	D	E*	E*	E*	E*	E*	W									
ldf Y(r1)➔f3					F	D	X	M	W											
ldf Y+4(r1)➔f7						F	D	X	M	s*	s*	W								
addf f2,f3➔f4							F	D	d*	E+	E+	s*	W							
addf f6,f7➔f8								F	p*	D	E+	p*	E+	W						
stf f4➔Z(r1)									F	D	X	M	W							
stf f8➔Z+4(r1)										F	D	X	M	W						
addi r1➔8,r1											F	D	X	M	W					
blt r1,r2,0												F	D	X	M	W				
ldf X(r1)➔f1													F	D	X	M	**W**			

+ Performance: 12 insn / 13 cycles = 0.92 IPC
+ Utilization: actual/peak IPC = 0.92 /1 = 92%
+ **Speedup**: (2 * 11 cycles) / 13 cycles = 1.69

17

Loop Unrolling Shortcomings

- Static code growth → more I$ misses (limits unrolling)
- Needs more registers to hold values (ISA limits this)
- Doesn't handle: non-loops, inter-iteration dependences

```
for (i=0;i<N;i++)
  X[i]=A*X[i-1];
```

```
ldf X-4(r1),f1                        ldf X-4(r1),f1
mulf f0,f1,f2                         mulf f0,f1,f2
stf f2,X(r1)                         stf f2,X(r1)
addi r1,4,r1                          mulf f0,f2,f3
blt r1,r2,0                           stf f3,X+4(r1)
ldf X-4(r1),f1                        addi r1,4,r1
mulf f0,f1,f2                         blt r1,r2,0
stf f2,X(r1)
addi r1,4,r1
blt r1,r2,0
```

- Two **mulf**'s are not parallel
- Other (more advanced) techniques help

18

Summary: Static Scheduling Limitations

- Limited number of registers (set by ISA)

- Scheduling scope
  - Example: hard to move memory insns past branches

- Inexact memory aliasing information
  - Often prevents reordering of loads above stores

- Caches misses (or any runtime event) confound scheduling
  - How can the compiler know which loads will miss/hit?
  - Can impact the compiler's scheduling decisions

Scheduling Techniques

- **Stall Removal**
  - Separate load-use pairs

- **Scope enlarging**
  - For Loops: **loop unrolling**
  - **For Non-loops:**
    - **Superblocks** (biased branches)
    - **Predication** (non-biased branches)

- **Exploit Data-Level Parallelism**
  - Vectors


Superblocks

Source code

```
A = Y[i];
if (A == 0)
    A = W[i];
else
    Y[i] = 0;
Z[i] = A*X[i];
```

Machine code

```
0: ldf Y(r1),f2
1: fbne f2,4
2: ldf W(r1),f2
3: jump 5
4: stf f0,Y(r1)
5: ldf X(r1),f4
6: mulf f4,f2,f6
7: stf f6,Z(r1)
```

4 basic blocks: A,B,C,D

A
```
0: ldf Y(r1),f2
1: fbne f2,4
```
NT=5%   T=95%

B
```
2: ldf W(r1),f2
3: jump 5
```
C
```
4: stf f0,Y(r1)
```

D
```
5: ldf X(r1),f4
6: mulf f4,f2,f6
7: stf f6,Z(r1)
```

- Use when branch is highly biased
- Fuse blocks of most frequent path: ACD
- Schedule
- Create **repair code** in case path = ABD

Superblock and Repair Code

Superblock
```
0: ldf Y(r1),f2
1: fbeq f2,2
4: stf f0,Y(r1)
5: ldf X(r1),f4
6: mulf f4,f2,f6
7: stf f6,Z(r1)
```

Repair code
```
2: ldf W(r1),f2
5': ldf X(r1),f4
6': mulf f4,f2,f6
7': stf f6,Z(r1)
```

- What did we do?
  - Change sense (test) of branch 1
    - Original taken target now fall-thru
  - Created repair block
    - May need to duplicate some code (here basic-block D)
  - Haven't actually scheduled superblock yet

Superblocks Scheduling I

Superblock
```
0: ldf Y(r1),f2
1: fbeq f2,2
5: ldf X(r1),f4
6: mulf f4,f2,f6
4: stf f0,Y(r1)
7: stf f6,Z(r1)
```

Repair code
```
2: ldf W(r1),f2
5': ldf X(r1),f4
6': mulf f4,f2,f6
7': stf f6,Z(r1)
```

- First scheduling move: move insns 5 and 6 above insn 4
  - Hmmm: moved load (5) above store (4)
  - We can tell this is OK, but can the compiler
    - If yes, fine
    - Otherwise, need to do something

Predication

- Conventional control
  - Conditionally executed insns also conditionally fetched
- **Predication**
  - Conditionally executed insns unconditionally fetched
  - **Full predication** (ARM, IA-64)
    - Tag every insn with predicate, costs extra bits
  - **Conditional moves** (Alpha, IA-32)
    - Construct appearance of full predication from one primitive
      `cmoveq r1,r2,r3        // if (r1==0) r3=r2;`
    - May require some code duplication to achieve desired effect
    + Only good way of adding predication to an existing ISA
- **If-conversion**: replacing control with predication
  + Good if branch is unpredictable (save mis-prediction)
  - But more instructions fetched and "executed"

Avoiding Branches via ISA: Predication

- Conventional control
  - Conditionally executed insns also conditionally fetched

	1	2	3	4	5	6	7	8	9
beq r3,targ	F	D	X	M	W				
sub r6,1,r5		F	D	--	--	--	flushed: wrong path		
targ:add r4,r5,r4			F	--	--	--	--	flushed: why?	
targ:add r4,r5,r4				F	D	X	M	W	

  - If **beq** mis-predicts, both **sub** and **add** must be flushed
  – Waste: **add** is independent of mis-prediction
- **Predication**: not prediction, predic**a**tion
  - ISA support for conditionally-executed unconditionally-fetched insns
  - If **beq** mis-predicts, annul **sub** in place, preserve **add**
    - Example is if-then, but if-then-else can be predicated too
  - How is this done? How does **add** get correct value for **r5**

25

Full Predication

Full predication
- Every insn can be annulled, annulment controlled by...
- Predicate registers: additional register in each insn (e.g., IA64)

	1	2	3	4	5	6	7	8	9
setp.eq r3,p3	F	D	X	M	W				
sub.p r6,1,r5,p3		F	D	X	--	--	annulled		
targ:add r4,r5,r4			F	D	X	M	W		

- Predicate codes: condition bits in each insn (e.g., ARM)

	1	2	3	4	5	6	7	8	9
setcc r3	F	D	X	M	W				
sub.nz r6,1,r5		F	D	X	--	--	annulled		
targ:add r4,r5,r4			F	D	X	M	W		

- Only ALU insn shown (**sub**), but this applies to all insns, even stores
- Branches replaced with "set-predicate" insns

26

Conditional Register Moves (CMOVs)

Conditional (register) moves
- Construct appearance of full predication from one primitive
  ```
  cmoveq r1,r2,r3        // if (r1==0) r3=r2;
  ```
- May require some code duplication to achieve desired effect
- Painful, potentially impossible for some insn sequences
- Requires more registers
- Only good way of retro-fitting predication onto ISA
  (e.g., IA32, Alpha)

	1	2	3	4	5	6	7	8	9
sub r6,1,r9		D	X	M	W				
cmovne r3,r9,r5		F	D	X	M	W			
targ:add r4,r5,r4			F	D	X	M	W		

27

Non-Biased Branches: Use Predication


```
A  0: ldf Y(r1),f2
   1: fbne f2,4
        NT=50%        T=50%
B  2: ldf W(r1),f2    4: stf f0,Y(r1)  C
   3: jump 5
D  5: ldf X(r1),f4
   6: mulf f4,f2,f6
   7: stf f6,Z(r1)
```

Using Predication
```
0: ldf Y(r1),f2
1: fspne f2,p1
2: ldf.p p1,W(r1),f2
4: stf.np p1,f0,Y(r1)
5: ldf X(r1),f4
6: mulf f4,f2,f6
7: stf f6,Z(r1)
```

28

ISA Support for Predication

```
0: ldf Y(r1),f2
1: fspne f2,p1
2: ldf.p p1,W(r1),f2
4: stf.np p1,f0,Y(r1)
5: ldf X(r1),f4
6: mulf f4,f2,f6
7: stf f6,Z(r1)
```

- IA-64: change branch 1 to **set-predicate insn fspne**
- Change insns 2 and 4 to **predicated insns**
  - **ldf.p** performs **ldf** if predicate **p1** is true
  - **stf.np** performs **stf** if predicate **p1** is false

29

Predication Performance

- Cost/benefit analysis
  - **Benefit**: predication avoids branches
    - Thus avoiding mis-predictions
    - Also reduces pressure on predictor table (few branches to track)
  - **Cost**: extra (annulled) instructions
- Since branch predictors are highly accurate...
  - Might not help:
    - 5-stage pipeline, two instruction on each path of if-then-else
    - No performance gain, likely slower if branch predictable
  - Or even hurt!
  - But can help:
    - Deeper pipelines, hard-to-predict branches, and few added insns
- Predication is useful, but not a panacea

30

Aside: Profiling – start here

How do we know whether a branch is biased or not?
Profile: statistical information about program tendencies
- Collect from previous program runs (different inputs)
± Works OK depending on information
  - Memory latencies (cache misses)
    + Which loads miss frequently independent of inputs?
    – Depends on cache configuration
  - Memory dependences
    - Which loads & stores communicate with each other?
    + Stable across inputs
  - Branch outcomes
    - Which branches are usually taken/not-taken?
    – Not so stable across inputs
- Popular research topic

Scheduling Techniques

- **Stall Removal**
  - Separate load-use pairs

- **Scope enlarging**
  - For Loops: **loop unrolling**
  - **For Non-loops:**
    - **Superblocks** (biased branches)
    - **Predication** (non-biased branches)

- **Exploit Data-Level Parallelism**
  - Vectors

Data-Level Parallelism

Data-level parallelism (DLP)
- Single operation repeated on multiple data elements
  - SIMD (**S**ingle-**I**nstruction, **M**ultiple-**D**ata)
- Less general than ILP: parallel insns are all same operation
- Exploit with **vectors**

Old idea: Cray-1 supercomputer from late 1970s
- Eight 64-entry x 64-bit floating point "Vector registers"
  - 4096 bits (0.5KB) in each register! 4KB vector register file
- Special vector instructions to perform vector operations
  - Load vector, store vector (wide memory operation)
  - Vector+Vector addition, subtraction, multiply, etc.
  - Vector+Constant addition, subtraction, multiply, etc.
  - In Cray-1, each instruction specifies 64 operations!

Example Vector ISA Extensions

Extend ISA with floating point (FP) vector storage …
- **Vector register**: fixed-size array of 32- or 64- bit FP elements
- **Vector length**: For example: 4, 8, 16, 64, …

- … and example operations for vector length of 4
  - Load vector: `ldf.v X(r1),v1 =`
    ```
    ldf X+0(r1),v1[0]
    ldf X+1(r1),v1[1]
    ldf X+2(r1),v1[2]
    ldf X+3(r1),v1[3]
    ```
  - Add two vectors: `addf.vv v1,v2,v3`
    `addf v1[i],v2[i],v3[i] (where i is 0,1,2,3)`
  - Add vector to scalar: `addf.vs v1,f2,v3`
    `addf v1[i],f2,v3[i]  (where i is 0,1,2,3)`

Example Use of Vectors – 4-wide

```
          ldf X(r1),f1                    ldf.v X(r1),v1
          mulf f0,f1,f2                   mulf.vs v1,f0,v2
          ldf Y(r1),f3                    ldf.v Y(r1),v3
          addf f2,f3,f4                   addf.vv v2,v3,v4
          stf f4,Z(r1)                    stf.v v4,Z(r1)
7x1024    addi r1,4,r1     7x256 insns    addi r1,16,r1
insns     blti r1,4096,0   (4x fewer insns) blti r1,4096,0
```

Operations
- Load vector: `ldf.v X(r1),v1`
- Multiply vector to scalar: `mulf.vs v1,f2,v3`
- Add two vectors: `addf.vv v1,v2,v3`
- Store vector: `stf.v v1,X(r1)`

Performance?
- If CPI = 1, 4x speedup
- CPI not always 1
  - Execution width (implementation) ≠ vector width (ISA)

Why Vectorization is Awesome

Have your cake and eat it, too
All the benefits of a wider machine, without superscalar costs
- Single instruction fetch
- Wide reads & writes (without multiple $ or regfile ports)
- Wider data to bypass ≠ N^2 bypass

Execution width (implementation) vs vector width (ISA)
- Example: Pentium 4 and Core 1 execute vector ops at half width
- Core 2 executes them at full width
- Intel's Sandy Bridge brings 256-bit vectors to x86
- Intel's Larrabee graphics chip brings 512-bit vectors to x86

Vector + superscalar? Sure!
- Multiple n-wide vector instructions per cycle

Scheduling: Compiler or Hardware

Compiler
+ Large scheduling scope (full program)
+ Simple hardware → fast clock, short pipeline, and low power
− Low branch prediction accuracy (profiling?)
− Little information on memory dependences (profiling?)
− Can't dynamically respond to cache misses (or anything really)
− Hard to speculate, recover from mis-speculation (h/w support?)

Hardware
− Finite buffering resources fundamentally limit scheduling scope
− Scheduling machinery adds pipeline stages and consumes power
+ High branch prediction accuracy
+ Dynamic information about memory dependences
+ Can respond to cache misses
+ Easy to speculate and recover from mis-speculation

37